

PostgreSQL Assignment-3

BY Manish Yadav

1. Create a new PostgreSQL database named "keenable".

```
postgres=# create database keenable;
```

Here's a short explanation of the command:

create database: This is the SQL command for creating a new database.

keenable: This is the name of the database being created. You can replace "keenable" with the desired name for your database.

This command is essential for initializing a new database within your PostgreSQL server.

```
postgres=# create database keenable;
CREATE DATABASE
postgres=#
postgres=#
postgres=# \l
```

Name	Owner	Encoding	Locale Provider	Collate	Ctype	ICU Locale	ICU Rules	Access privileges
keenable	postgres	UTF8	libc	en_IN	en_IN			
manish_rao	postgres	UTF8	libc	en_IN	en_IN			
manishapi	postgres	UTF8	libc	en_IN	en_IN			
manishdb	postgres	UTF8	libc	en_IN	en_IN			=Tc/postgres + postgres=CTc/postgres+ manish=CTc/postgres
my_first_pgdb	postgres	UTF8	libc	en_IN	en_IN			
postgres	postgres	UTF8	libc	en_IN	en_IN			
raoofindia	postgres	UTF8	libc	en_IN	en_IN			
template0	postgres	UTF8	libc	en_IN	en_IN			=c/postgres + postgres=CTc/postgres
template1	postgres	UTF8	libc	en_IN	en_IN			=c/postgres + postgres=CTc/postgres

(9 rows)

2. Define three schemas within the database: "hr", "technical", and "management".

```
postgres=# \c keenable
```

```
postgres=# \c keenable
psql (16.0 (Ubuntu 16.0-1.pgdg22.04+1), server 14.9 (Ubuntu 14.9-1.pgdg22.04+1))
You are now connected to database "keenable" as user "postgres".
```

It's a way to change current database context to "keenable" for running SQL queries and operations within that database.

- A. `keenable=# create schema hr;`
- B. `keenable=# create schema technical;`
- C. `keenable=# create schema management;`

```

keenable=# create schema hr;
CREATE SCHEMA
keenable=#
keenable=#
keenable=# create schema technical;
CREATE SCHEMA
keenable=# create schema management;
CREATE SCHEMA

```

This query creates a schema named "hr","technical","management" in the "keenable" database.

keenable=# \dn

```

keenable=# \dn
      List of schemas
  Name      | Owner
-----+-----
 hr         | postgres
 management | postgres
 public     | postgres
 technical  | postgres
(4 rows)

```

The command \dn in PostgreSQL is used to list all the schemas in the current database.

2.1 HR Schema:

```

keenable=# CREATE TABLE hr.employees (
    employee_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    email VARCHAR(100),
    hire_date DATE
);
CREATE TABLE

keenable=# CREATE TABLE hr.vacation_requests (
    request_id SERIAL PRIMARY KEY,
    employee_id INT,
    start_date DATE,
    end_date DATE,
    status VARCHAR(20),
    FOREIGN KEY (employee_id) REFERENCES hr.employees(employee_id)
);
CREATE TABLE

```

Create the following tables in the "hr" schema:

2.1.1.

```

keenable=#
keenable=#
INSERT INTO hr.employees (first_name, last_name, email, hire_date)
VALUES
 ('Rahul', 'Gupta', 'rahul.gupta@example.com', '1990-05-15'),
 ('Priya', 'Sharma', 'priya.sharma@example.com', '1995-08-20'),
 ('Amit', 'Patel', 'amit.patel@example.com', '1987-03-10'),
 ('Sneha', 'Singh', 'sneha.singh@example.com', '2000-11-28'),
 ('Deepak', 'Verma', 'deepak.verma@example.com', '1992-07-05');
INSERT 0 5
keenable=#
keenable=# \dt
Did not find any relations.
keenable=# \dt hr.employees ;
      List of relations
 Schema | Name      | Type  | Owner
-----+-----+-----+-----
 hr     | employees | table | postgres
(1 row)

\dt: extra argument ";" ignored
keenable=# select * from hr.employees;
 employee_id | first_name | last_name | email | hire_date
-----+-----+-----+-----+-----
          1 | Rahul     | Gupta    | rahul.gupta@example.com | 1990-05-15
          2 | Priya     | Sharma   | priya.sharma@example.com | 1995-08-20
          3 | Amit      | Patel    | amit.patel@example.com   | 1987-03-10
          4 | Sneha     | Singh    | sneha.singh@example.com  | 2000-11-28
          5 | Deepak    | Verma    | deepak.verma@example.com | 1992-07-05
(5 rows)
keenable=#

```

"employees" to store employee information (employee_id, first_name,last_name, email, hire_date).

2.1.2.

```
keenable=# INSERT INTO hr.vacation_requests (employee_id, start_date, end_date, status)
VALUES
(1, '2023-01-10', '2023-01-15', 'Approved'),
(2, '2023-03-20', '2023-03-25', 'Pending'),
(3, '2023-02-05', '2023-02-12', 'Approved'),
(4, '2023-04-10', '2023-04-15', 'Rejected'),
(5, '2023-05-01', '2023-05-10', 'Pending');
INSERT 0 5
```

```
keenable=# select * from hr.vacation_requests;
 request_id | employee_id | start_date | end_date | status
-----+-----+-----+-----+-----
          1 |           1 | 2023-01-10 | 2023-01-15 | Approved
          2 |           2 | 2023-03-20 | 2023-03-25 | Pending
          3 |           3 | 2023-02-05 | 2023-02-12 | Approved
          4 |           4 | 2023-04-10 | 2023-04-15 | Rejected
          5 |           5 | 2023-05-01 | 2023-05-10 | Pending
(5 rows)
```

"vacation_requests" to track employee vacation requests (request_id,employee_id, start_date, end_date, status).

2.2. Technical Schema:

```
keenable=# CREATE TABLE technical.projects (
    project_id SERIAL PRIMARY KEY,
    project_name VARCHAR(100),
    start_date DATE,
    end_date DATE
);
CREATE TABLE
keenable=# CREATE TABLE technical.project_assignments (
    assignment_id SERIAL PRIMARY KEY,
    project_id INT,
    employee_id INT,
    assignment_date DATE,
    FOREIGN KEY (project_id) REFERENCES technical.projects(project_id),
    FOREIGN KEY (employee_id) REFERENCES hr.employees(employee_id)
);
CREATE TABLE
```

Create the following tables in the "technical" schema:

2.2.1.

```
keenable=# INSERT INTO technical.projects (project_name, start_date, end_date)
VALUES
('Project A', '2023-01-10', '2023-02-15'),
('Project B', '2023-03-20', '2023-04-25'),
('Project C', '2023-02-05', '2023-03-12'),
('Project D', '2023-04-10', '2023-05-15'),
('Project E', '2023-05-01', '2023-06-10');
INSERT 0 5
keenable=#
```

```
keenable=# select * from technical.projects;
 project_id | project_name | start_date | end_date
-----+-----+-----+-----
          1 | Project A    | 2023-01-10 | 2023-02-15
          2 | Project B    | 2023-03-20 | 2023-04-25
          3 | Project C    | 2023-02-05 | 2023-03-12
          4 | Project D    | 2023-04-10 | 2023-05-15
          5 | Project E    | 2023-05-01 | 2023-06-10
(5 rows)

keenable=#
```

"projects" to store project details (project_id, project_name, start_date,end_date).

2.2.2.

```

keenable=#
keenable=# INSERT INTO technical.project_assignments (project_id, employee_id, assignment_date)
VALUES
  (1, 1, '2023-01-15'),
  (2, 2, '2023-03-20'),
  (3, 3, '2023-02-07'),
  (4, 4, '2023-04-12'),
  (5, 5, '2023-05-03');
INSERT 0 5
keenable=#

```

```

keenable=# select * from technical.project_assignments;
assignment_id | project_id | employee_id | assignment_date
-----+-----+-----+-----
          1 |          1 |           1 | 2023-01-15
          2 |          2 |           2 | 2023-03-20
          3 |          3 |           3 | 2023-02-07
          4 |          4 |           4 | 2023-04-12
          5 |          5 |           5 | 2023-05-03
(5 rows)
keenable=#

```

"project_assignments" to associate employees with projects (assignment_id, project_id, employee_id, assignment_date).

2.3. Management Schema:

```

keenable=# CREATE TABLE management.departments (
    department_id SERIAL PRIMARY KEY,
    department_name VARCHAR(100),
    location VARCHAR(100)
);
CREATE TABLE

keenable=# CREATE TABLE management.departments_employees (
    assignment_id SERIAL PRIMARY KEY,
    employee_id INT,
    department_id INT,
    start_date DATE,
    end_date DATE,
    FOREIGN KEY (employee_id) REFERENCES hr.employees(employee_id),
    FOREIGN KEY (department_id) REFERENCES management.departments(department_id)
);
CREATE TABLE
keenable=#
keenable=#

```

Create the following tables in the "management" schema:

2.3.1.

```

keenable=# INSERT INTO management.departments (department_name, location)
VALUES
  ('HR Department', 'Mumbai'),
  ('IT Department', 'Bangalore'),
  ('Finance Department', 'Delhi'),
  ('Marketing Department', 'Chennai'),
  ('Operations Department', 'Kolkata');
INSERT 0 5
keenable=# \do

```

```

keenable=# select * from management.departments;
department_id | department_name | location
-----+-----+-----
          1 | HR Department   | Mumbai
          2 | IT Department   | Bangalore
          3 | Finance Department | Delhi
          4 | Marketing Department | Chennai
          5 | Operations Department | Kolkata
(5 rows)

```

"departments" to manage department information (department_id, department_name, location).

2.3.2.

```

keenable=# CREATE TABLE management.departments_employees (
  assignment_id SERIAL PRIMARY KEY,
  employee_id INT,
  department_id INT,
  start_date DATE,
  end_date DATE);
CREATE TABLE

keenable=#
keenable=# select * from management.departments_employees;
 assignment_id | employee_id | department_id | start_date | end_date
-----
1 | 1 | 1 | 2023-01-15 | 2023-04-12
2 | 2 | 2 | 2023-03-20 | 2023-05-31
3 | 3 | 3 | 2023-02-07 | 2023-06-30
4 | 4 | 4 | 2023-04-12 | 2023-07-15
5 | 5 | 5 | 2023-05-03 | 2023-08-15
(5 rows)

keenable=#

```

"department_employees" to track employee department assignments (assignment_id, employee_id, department_id, start_date, end_date).

3. Relationships:

Establish appropriate foreign key relationships between tables to maintain referential integrity. For example, link "project_assignments" to "projects" and "employees."

For that I had used the following command:

1. keenable=#-- In the "technical" schema ALTER TABLE technical.project_assignments ADD CONSTRAINT fk_project_assignment_project FOREIGN KEY (project_id) REFERENCES technical.projects(project_id); ALTER TABLE
2. keenable=# -- In the "technical" schema ALTER TABLE technical.project_assignments ADD CONSTRAINT fk_project_assignment_employee FOREIGN KEY (employee_id) REFERENCES hr.employees(employee_id); ALTER TABLE

```

keenable=# -- In the "technical" schema
ALTER TABLE technical.project_assignments
ADD CONSTRAINT fk_project_assignment_project
FOREIGN KEY (project_id) REFERENCES technical.projects(project_id);
ALTER TABLE
keenable=# -- Show the details of the foreign key constraint
\d technical.project_assignments

```

Column	Type	Collation	Nullable	Default
assignment_id	integer		not null	nextval('technical.project_assignments_assignment_id_seq'::regclass)
project_id	integer			
employee_id	integer			
assignment_date	date			

```

Indexes:
    "project_assignments_pkey" PRIMARY KEY, btree (assignment_id)
Foreign-key constraints:
    "fk_project_assignment_project" FOREIGN KEY (project_id) REFERENCES technical.projects(project_id)
    "project_assignments_employee_id_fkey" FOREIGN KEY (employee_id) REFERENCES hr.employees(employee_id)
    "project_assignments_project_id_fkey" FOREIGN KEY (project_id) REFERENCES technical.projects(project_id)

```

```

keenable=# -- In the "technical" schema
ALTER TABLE technical.project_assignments
ADD CONSTRAINT fk_project_assignment_employee
FOREIGN KEY (employee_id) REFERENCES hr.employees(employee_id);
ALTER TABLE
keenable=# -- Show the details of the foreign key constraint
\d technical.project_assignments

```

Column	Type	Collation	Nullable	Default
assignment_id	integer		not null	nextval('technical.project_assignments_assignment_id_seq'::regclass)
project_id	integer			
employee_id	integer			
assignment_date	date			

```

Indexes:
    "project_assignments_pkey" PRIMARY KEY, btree (assignment_id)
Foreign-key constraints:
    "fk_project_assignment_employee" FOREIGN KEY (employee_id) REFERENCES hr.employees(employee_id)
    "fk_project_assignment_project" FOREIGN KEY (project_id) REFERENCES technical.projects(project_id)
    "project_assignments_employee_id_fkey" FOREIGN KEY (employee_id) REFERENCES hr.employees(employee_id)
    "project_assignments_project_id_fkey" FOREIGN KEY (project_id) REFERENCES technical.projects(project_id)

```

```

keenable=#

```

4. Queries and Reporting:

Write SQL queries to retrieve information such as: 4.1. Employees in a specific department 4.2. Projects assigned to an employee. 4.3. Vacation requests for an employee.

4.1.

For that I had used the following command:

```

keenable=# SELECT e.first_name, e.last_name FROM hr.employees e JOIN
management.departments_employees de ON e.employee_id = de.employee_id JOIN
management.departments d ON de.department_id = d.department_id WHERE d.department_name = 'HR
Department';

```

```

keenable=# SELECT e.first_name, e.last_name
FROM hr.employees e
JOIN management.departments_employees de ON e.employee_id = de.employee_id
JOIN management.departments d ON de.department_id = d.department_id
WHERE d.department_name = 'HR Department';

```

first_name	last_name
Rahul	Gupta

```

(1 row)

```

Employees in a specific department

4.2.

For that I had used the following command:

keenable=# SELECT p.project_name FROM technical.project_assignments pa INNER JOIN technical.projects p ON pa.project_id = p.project_id WHERE pa.employee_id = 1; -- Replace with the desired employee ID

```
keenable=# SELECT p.project_name
FROM technical.project_assignments pa
INNER JOIN technical.projects p ON pa.project_id = p.project_id
WHERE pa.employee_id = 1; -- Replace with the desired employee ID
 project_name
-----
Project A
(1 row)
```

Projects assigned to an employee.

4.3.

For that I had used the following command:

keenable=# SELECT vr.request_id, vr.start_date, vr.end_date, vr.status FROM hr.vacation_requests vr WHERE vr.employee_id = 4; -- Replace with the desired employee ID

```
keenable=# SELECT vr.request_id, vr.start_date, vr.end_date, vr.status
FROM hr.vacation_requests vr
WHERE vr.employee_id = 4; -- Replace with the desired employee ID
 request_id | start_date | end_date | status
-----+-----+-----+-----
          4 | 2023-04-10 | 2023-04-15 | Rejected
(1 row)

keenable=#
```

Vacation requests for an employee.